

Big Data And Hadoop

Session 16 – Assignment 3

Problem Statement:


Link: <https://acadgild.com/blog/transactions-in-hive/>

Refer the above given link for transactions in Hive and implement the operations given in the blog using your own sample data set and send us the screenshot.

Solution:

1. Row Level Transactions:

Transactions are provided at the row-level in Hive 0.14. So we verify the version of hive installed using the **version** option as follows:



```
acadmild@localhost:~  
File Edit View Search Terminal Help  
[acadmild@localhost ~]$ hive --version  
Hive 0.14.0  
Subversion file:///Users/ghagleitner/Projects/hive-svn/rel-prep/hive-14-rel-prep -r Unknown  
Compiled by ghagleitner on Sat Nov 8 23:25:06 PST 2014  
From source with checksum 49c2182a0856f7917f571802a7594b00  
[acadmild@localhost ~]$
```

The different row-level transactions available in Hive 0.14 are as follows:

1. Insert
2. Delete
3. Update

We start the hive command line by executing the command `hive` as shown below:



```
acadmild@localhost:~  
File Edit View Search Terminal Help  
[acadmild@localhost ~]$ hive  
Logging initialized using configuration in jar:file:/usr/local/hive/lib/hive-common-0.14.0.jar!/hive-log4j.properties  
SLF4J: Class path contains multiple SLF4J bindings.  
SLF4J: Found binding in [jar:file:/usr/local/hive/lib/hive-jdbc-0.14.0-standalone.jar!/org/slf4j/impl/StaticLoggerBinder.class]  
SLF4J: Found binding in [jar:file:/usr/local/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]  
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.  
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]  
hive>
```

The above snapshot also shows that hive prompt has started. A pre-requisite to use hive is to start mysql server. This was done using the command `sudo service mysqld start`.

Before creating a Hive table that supports transactions, the transaction features present in Hive needs to be turned on, as by default they are turned off.

The below properties needs to be set appropriately in *hive shell*, order-wise to work with transactions in Hive:

- `Hive.support.concurrency = true;`
- `Hive.enforce.bucketing = true;`
- `Hive.exec.dynamic.partition.mode = nonstrict;`
- `Hive.txn.manager = org.apache.hadoop.hive.ql.lockmgr.DbTxnManager;`
- `Hive.compactor.initiator.on = true;`

- Hive.compactor.worker.threads = 1;

These properties are set as follows:

```

acadgild@localhost:~
File Edit View Search Terminal Help

hive> set hive.support.concurrency = true;
hive> set hive.enforce.bucketing = true;
hive> set hive.exec.dynamic.partition.mode = nonstrict;
hive> set hive.txn.manager = org.apache.hadoop.hive.ql.lockmgr.DbTxnManager;
hive> set hive.compactor.initiator.on = true;
hive> set hive.compactor.worker.threads = 1;
hive>

```

2. Create Table that supports hive transactions:

We create a table with name 'employee' and the columns present in the table are *id*, *name*, *location*, *salary*. We are *bucketing* the table by 'id' and the table format is 'orc', also we are enabling the transactions in the table by specifying it inside the *TBLPROPERTIES* as 'transactional'='true'.

The query used is

CREATE TABLE employee(id int, name string, location string, salary int) clustered by (id) into 2 buckets stored as orc TBLPROPERTIES("transactional"= "true");

```

acadgild@localhost:~
File Edit View Search Terminal Help

hive> CREATE TABLE employee(id int, name string, location string, salary int) clustered by (id) into 2 buckets stored as orc
TBLPROPERTIES("transactional"="true");
OK
Time taken: 0.439 seconds
hive>

```

We can see the table created using the command **SHOW TABLES:**

```

acadgild@localhost:~
File Edit View Search Terminal Help

hive> SHOW TABLES;
OK
employee
first
locations
toss
users
Time taken: 0.087 seconds, Fetched: 5 row(s)
hive>

```

3. Inserting data into hive table:

The command is used to insert row wise data into the Hive table is

INSERT INTO TABLE employee values (101, "Amit", "Pune", 23000), (102, "Akash", "Pune", 20000), (103, "Stefy", "Mumbai", 48000), (104, "Sanket", "Mumbai", 41000), (105, "Abil", "Pune", 50000);

Here, each row is separated by ' () ' brackets.

```
acadgild@localhost:~  
File Edit View Search Terminal Help  
hive> INSERT INTO TABLE employee values (101, "Amit", "Pune", 23000), (102, "Akash", "Pune", 20000), (103, "Stefy", "Mumbai",  
48000), (104, "Sanket", "Mumbai", 41000), (105, "Abil", "Pune", 50000);  
Query ID = acadgild_20171017170202_05cf30d4-54bb-468b-a525-e2ba022af011  
Total jobs = 1  
Launching Job 1 out of 1  
Number of reduce tasks determined at compile time: 2  
In order to change the average load for a reducer (in bytes):  
  set hive.exec.reducers.bytes.per.reducer=<number>  
In order to limit the maximum number of reducers:  
  set hive.exec.reducers.max=<number>  
In order to set a constant number of reducers:  
  set mapreduce.job.reduces=<number>  
Starting Job = job_1508233734419_0001, Tracking URL = http://localhost:8088/proxy/application_1508233734419_0001/  
Kill Command = /home/acadgild/hadoop-2.6.0/bin/hadoop job -kill job_1508233734419_0001  
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 2  
2017-10-17 17:03:08,933 Stage-1 map = 0%, reduce = 0%  
2017-10-17 17:03:18,060 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.9 sec  
2017-10-17 17:03:33,199 Stage-1 map = 100%, reduce = 50%, Cumulative CPU 6.96 sec  
2017-10-17 17:03:35,397 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 12.69 sec  
MapReduce Total cumulative CPU time: 12 seconds 690 msec  
Ended Job = job_1508233734419_0001  
Loading data to table default.employee  
Table default.employee stats: [numFiles=2, numRows=5, totalSize=1025, rawDataSize=924]  
MapReduce Jobs Launched:  
Stage-Stage-1: Map: 1 Reduce: 2 Cumulative CPU: 12.69 sec HDFS Read: 384 HDFS Write: 1171 SUCCESS  
Total MapReduce CPU Time Spent: 12 seconds 690 msec  
OK  
Time taken: 46.562 seconds  
hive>
```

The contents of the table can be viewed using the command **select * from employee;**

```
acadgild@localhost:~  
File Edit View Search Terminal Help  
hive> select * from employee;  
OK  
104      Sanket      Mumbai      41000  
102      Akash       Pune        20000  
105      Abil        Pune        50000  
103      Stefy       Mumbai      48000  
101      Amit        Pune        23000  
Time taken: 0.18 seconds, Fetched: 5 row(s)  
hive>
```

Now if we try to re-insert the same data again, it will be appended to the previous data as shown below:

```
acadgild@localhost:~  
File Edit View Search Terminal Help  
hive> INSERT INTO TABLE employee values (101, "Amit", "Pune", 23000), (102, "Akash", "Pune", 20000), (103, "Stefy", "Mumbai",  
48000), (104, "Sanket", "Mumbai", 41000), (105, "Abil", "Pune", 50000);  
Query ID = acadgild_20171017170505_1533b156-082f-40c9-9404-3fde48a265f7  
Total jobs = 1  
Launching Job 1 out of 1  
Number of reduce tasks determined at compile time: 2  
In order to change the average load for a reducer (in bytes):  
  set hive.exec.reducers.bytes.per.reducer=<number>  
In order to limit the maximum number of reducers:  
  set hive.exec.reducers.max=<number>  
In order to set a constant number of reducers:  
  set mapreduce.job.reduces=<number>  
Starting Job = job_1508233734419_0002, Tracking URL = http://localhost:8088/proxy/application_1508233734419_0002/  
Kill Command = /home/acadgild/hadoop-2.6.0/bin/hadoop job -kill job_1508233734419_0002  
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 2  
2017-10-17 17:06:07,753 Stage-1 map = 0%, reduce = 0%  
2017-10-17 17:06:15,382 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.29 sec  
2017-10-17 17:06:30,002 Stage-1 map = 100%, reduce = 50%, Cumulative CPU 6.4 sec  
2017-10-17 17:06:31,064 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 10.82 sec  
MapReduce Total cumulative CPU time: 10 seconds 820 msec  
Ended Job = job_1508233734419_0002  
Loading data to table default.employee  
Table default.employee stats: [numFiles=4, numRows=10, totalSize=2050, rawDataSize=1848]  
MapReduce Jobs Launched:  
Stage-Stage-1: Map: 1 Reduce: 2 Cumulative CPU: 10.82 sec HDFS Read: 384 HDFS Write: 1171 SUCCESS  
Total MapReduce CPU Time Spent: 10 seconds 820 msec  
OK  
Time taken: 36.003 seconds  
hive>
```

The repeated records seen using **SELECT * FROM employee;**

```
acadgild@localhost:~  
File Edit View Search Terminal Help  
hive> select * from employee;  
OK  
104 Sanket Mumbai 41000  
102 Akash Pune 20000  
104 Sanket Mumbai 41000  
102 Akash Pune 20000  
105 Abil Pune 50000  
103 Stefy Mumbai 48000  
101 Amit Pune 23000  
105 Abil Pune 50000  
103 Stefy Mumbai 48000  
101 Amit Pune 23000  
Time taken: 0.112 seconds, Fetched: 10 row(s)  
hive>
```

4. Updating the data in the hive table:

We can update the records using the update command as follows:

UPDATE employee SET name = 'Diksha' WHERE id = '105';

If we try to update the values of column used for bucketing, it throws an exception, this means that the Update command is not supported on the columns that are bucketed.

The updated data can be seen as follows:



```
acadgild@localhost:~  
File Edit View Search Terminal Help  
hive> SELECT * FROM employee;  
OK  
104 Sanket Mumbai 41000  
102 Akash Pune 20000  
104 Sanket Mumbai 41000  
102 Akash Pune 20000  
105 Diksha Pune 50000  
103 Stefy Mumbai 48000  
101 Amit Pune 23000  
105 Diksha Pune 50000  
103 Stefy Mumbai 48000  
101 Amit Pune 23000  
Time taken: 0.193 seconds, Fetched: 10 row(s)  
hive>
```

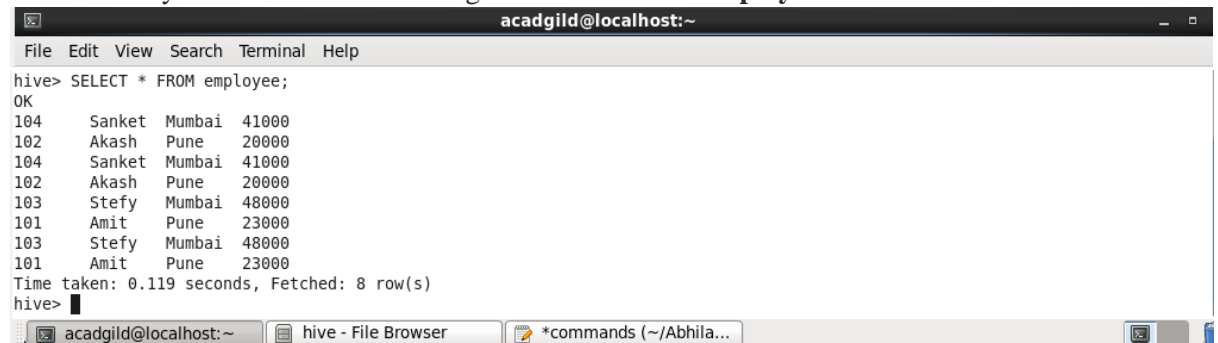
5. Deleting row from hive table:

The command used to delete the row is as follows

DELETE from employee where id = 105;

This command will delete the record which satisfies the condition **id = 105**.

We can verify the deletion of data using **SELECT * from employee**.



```
acadgild@localhost:~  
File Edit View Search Terminal Help  
hive> SELECT * FROM employee;  
OK  
104 Sanket Mumbai 41000  
102 Akash Pune 20000  
104 Sanket Mumbai 41000  
102 Akash Pune 20000  
103 Stefy Mumbai 48000  
101 Amit Pune 23000  
103 Stefy Mumbai 48000  
101 Amit Pune 23000  
Time taken: 0.119 seconds, Fetched: 8 row(s)  
hive>
```

From the output, we observe that the records with id 105 have been eliminated.